

# Predicting Household Behavior using Census Data

*Analysis of the accuracy, sensitivity, and interpretability of various classification algorithms in predicting household tenure (renting vs. owning) using Census PUMS data.*

*Civil & Environmental Engineering 263n: Scalable Spatial Analytics at UC-Berkeley Paul Sohn, December 11, 2016*

## Introduction

As a data scientist for UrbanSim, I am interested in methods to accurately and reliably predict household behavior. One of the household behaviors that cities and regions may be most interested in predicting is tenure: whether households choose to rent or own their units. This is one of several “choices” in the sequence of UrbanSim models where a method for classification is needed, such as the decision of what type of building to live in, what location to select.

Logit models (binary logit, multinomial logit) are the classification method of choice at this time. Given the much larger array of accessible classification algorithms that are available today, though, I am interested in exploring how the current implementation of logit models stacked in up performance against a range of other classifiers, all of which are available in the scikit-learn library. This project therefore aimed to evaluate selected classifiers from scikit-learn along the following dimensions:

- Accuracy: For what proportion of test cases does a trained classifier correctly predict tenure?
- Sensitivity: How much do changes in features (variables) affect classifier results?
- Interpretability: How easy is it to interpret classifier parameters and why a classifier predicts a certain outcome for a given case?

## Data

For this project, I used Public Use Microdata Sample (PUMS) data from the U.S. Census Bureau. This is a publicly available, disaggregate data source that provides records for individual persons and housing units. Information about renting and owning is available in the housing unit records of PUMS; I used the 2015 one-year sample data for California for this project. There were 155,218 records to begin with. Data were cleaned to remove records for institutional or non-institutional group quarters, so as to focus only on housing units. The dataset was also filtered to include only housing units for which the occupant(s) moved in 12 months ago or more recently. Focusing on recent movers was important for this analysis, since these households are likeliest to have characteristics similar to when they actually chose to rent or own their housing unit. The two steps above yielded a dataset of 15,928 records.

There are more than 150 variables in the housing unit record, but I chose to focus on a small subset that seemed applicable to the tenure decision. These include variables related to the *housing unit* itself, like the number of rooms or access to the internet, as well as variables related to the *household*, such as income or number of vehicles. It is important to distinguish between these types because while housing characteristics are available in this dataset, we would not use them for predicting tenure decisions, since they are specific to a housing unit and not the broader decision of whether to rent or buy. In the theory of this model in the UrbanSim context, households choose their tenure before deciding on a location. Nevertheless, we keep both types of variables in order to test their performance. The full list of variables is presented in Table 1.

Table 1: Variables used from the PUMS Housing Unit record

Variable	Description	Type
ACCESS	Access to the Internet	Housing
BATH	Bathtub or shower	Housing
RMSP	Number of rooms	Housing

Variable	Description	Type
YBL	When structure first built	Housing
KIT	Complete kitchen facilities	Housing
FS	Yearly food stamp/Supplemental Nutrition Assistance Program reciprocity	Household
LAPTOP	Laptop, desktop, or notebook computer	Household
VEH	Vehicles (1 ton or less) available	Household
HHL	Household language	Household
HHT	Household/family type (e.g. married couple, nonfamily)	Household
HINCP	Household income (past 12 months)	Household
HUGCL	Household with grandparent living with grandchildren	Household
HUPAC	HH presence and age of children	Household
LNGI	Limited English speaking household	Household
MULTG	Multigenerational Household	Household
NR	Presence of nonrelative in household	Household
PARTNER	Unmarried partner household	Household
SSMC	Same-sex married couple households	Household

Many of these are numerically coded nominal variables, with integer values representing different categories; the household language variable is a good example, with “1” representing “English only,” “2” representing “Spanish,” etc. When using many methods such as linear or logistic regression, it is important to recode these types of variables to dummy variables or other types that have a meaningful numeric coding, while making sure to leave one category out. Other methods like decision trees handle categorical variables well and do not require such recoding. I created recoded versions of many of these variables, shown in Table 2.

Table 2: Recoded variables

Original variable	Recoded variable	Description of recoding
ACCESS	access_recode	Convert from [1,2] to [1,0] coding
BATH	bath_recode	Convert from [1,2] to [1,0] coding
YBL	before1990	Convert to dummy variable with 1 for units built before 1990
KIT	kit_recode	Convert from [1,2] to [1,0] coding
FS	fs_recode	Convert from [1,2] to [1,0] coding
LAPTOP	laptop_recode	Convert from [1,2] to [1,0] coding
HHL	english_hh	Convert to dummy variable with 1 for households that are English-speaking
HHT	single_parent	Convert to dummy variable with 1 for family households with a single parent
HHT	nonfamily	Convert to dummy variable with 1 for non-family households
HUPAC	children	Convert to dummy variable with 1 for households with children
LNGI	good_eng_speaker	Convert from [1,2] to [1,0] coding
MULTG	multigen	Convert from [1,2] to [1,0] coding
PARTNER	unmarried_partner	Convert from [1,2] to [1,0] coding
SSMC	samesex_marriage	Convert from [1,2] to [1,0] coding

Between variable types (housing vs. household) and recoding statuses (original vs. recoded), four separate datasets were created:

- Original variables: housing and household variables
- Original variables: household variables only
- Recoded variables: housing and household variables
- Recoded variables: household variables only

The `data_preparation.py` module performs each of the steps described above.

## Classifiers

I selected several classifiers available in the scikit-learn library to test against the logit classifier that the current UrbanSim stack uses. Instead of using the custom implementation of multinomial and binary logit that exists in the UrbanSim library, I used the scikit-learn logistic regression classifier for simplicity and comparability. Logistic regression and logit classification are essentially identical methodologies for the purposes of this evaluation.

The classifiers I selected for this project are:

- Decision trees (scikit-learn estimator `DecisionTreeClassifier`)
- Random forests (`RandomForestClassifier`)
- Linear support vector machine (`LinearSVC`)
- K-nearest neighbors (`KNeighborsClassifier`)
- Logistic regression (`LogisticRegression`)

## Hyperparameter Tuning

Because some classifiers can yield substantially different accuracy results with different values of hyperparameters, I performed a randomized grid search on selected parameters using the scikit-learn `RandomizedSearchCV` function, which tests random combinations of parameters and returns the sets with the highest cross-validation accuracy score. I used the household only recoded dataset for this grid search. Table 3 shows parameters I tested for and the final parameters found for each. Note that I used the same parameter ranges for Decision Tree and Random Forests classifiers. The final value was only slightly different between the two; Random Forests had a `max_features` value of 9; the Decision Tree classifier had a value of 7.

Table 3: Hyperparameter tuning

Classifier	Parameter	Distribution	Final Value
Tree Classifiers	<code>max_depth</code>	<code>randint(2, 10)</code>	7
Tree Classifiers	<code>max_features</code>	<code>randint(1, 11)</code>	9 / 7
Tree Classifiers	<code>min_samples_split</code>	<code>randint(1, 11)</code>	2
Tree Classifiers	<code>min_samples_leaf</code>	<code>randint(1, 11)</code>	2
Tree Classifiers	<code>criterion</code>	["gini", "entropy"]	gini
LinearSVC	<code>C</code>	<code>np.linspace(0.1, 2, 20)</code>	1.4
LinearSVC	<code>loss</code>	['hinge', 'squared_hinge']	hinge
KNeighborsClassifier	<code>n_neighbors</code>	<code>randint(1, 1000)</code>	813
KNeighborsClassifier	<code>weights</code>	['uniform', 'distance']	uniform
KNeighborsClassifier	<code>algorithm</code>	['ball_tree', 'kd_tree']	ball_tree
KNeighborsClassifier	<code>leaf_size</code>	<code>randint(10, 100)</code>	58
KNeighborsClassifier	<code>p</code>	[1, 2]	1

## Evaluation: Accuracy

The first and most important criterion for evaluating these classifiers is the accuracy of predictions. The specific metric I used is mean accuracy from five-fold cross validation, using the `cross_val_score` function from scikit-learn's metrics library. This built-in function takes scikit-learn estimators, along with a training dataset, and performs k-fold cross validation and returns a selected metric for all  $k$  folds.

I iterated through all four datasets above, along with tuned classifiers, and took the mean cross-validation score. Results are presented below in Table 3.

Table 4: Accuracy results for five classifiers on four datasets

Variable Set	Recoding Status	Classifier	Mean Accuracy
Housing and Household	Not recoded	RandomForestClassifier	79.44%
Housing and Household	Not recoded	DecisionTreeClassifier	78.26%
Housing and Household	Not recoded	LogisticRegressionCV	76.10%
Housing and Household	Not recoded	KNeighborsClassifier	74.76%
Housing and Household	Not recoded	LinearSVC	74.33%
Housing and Household	Recoded	RandomForestClassifier	79.27%
Housing and Household	Recoded	DecisionTreeClassifier	78.05%
Housing and Household	Recoded	KNeighborsClassifier	74.79%
Housing and Household	Recoded	LogisticRegressionCV	74.44%
Housing and Household	Recoded	LinearSVC	67.75%
Household only	Not recoded	RandomForestClassifier	76.79%
Household only	Not recoded	DecisionTreeClassifier	75.96%
Household only	Not recoded	LogisticRegressionCV	75.29%
Household only	Not recoded	KNeighborsClassifier	74.68%
Household only	Not recoded	LinearSVC	73.35%
Household only	Recoded	RandomForestClassifier	76.82%
Household only	Recoded	DecisionTreeClassifier	76.40%
Household only	Recoded	KNeighborsClassifier	74.77%
Household only	Recoded	LogisticRegressionCV	74.42%
Household only	Recoded	LinearSVC	71.21%

The best results across all datasets are produced by the Random Forests estimator, with the Decision Tree classifier coming in second for each dataset. It's not surprising that the ensemble method outperforms the basic decision tree estimator, but it is interesting to note that tree-based methods are most accurate. Logistic Regression and K Nearest Neighbors Classifier have similar results, and Linear SVC comes in last place on each dataset. Note that I only tested the linear kernel of the support vector machine classifier; non-linear methods were not tested because fitting those algorithms is typically very slow relative to other methods. For the rest of this report, I will not evaluate the SVC classifier.

## Evaluation: Sensitivity

It is also helpful to better understand how sensitive each of the models are to various features. For example, how sensitive is the rent vs. own decision to household income in each of the models? For this evaluation, I made use of the open source ML Insights library, which provides functionality "to see how [a given] model performs when one feature is changed, holding everything else constant." I was hoping to see two things in each model: first, that it was sensitive to changes in important variables, particularly income and vehicle

ownership, and second, that there was consistency in sensitivity; in other words, if a model is sensitive to a variable, it is similarly sensitive across many changes.

I created a training set with 70% of the data from the household only recoded dataset, and tested sensitivity using the 30% testing dataset. The primary tool I used for interpreting sensitivity was the Feature Effect Summary method from the ML Insights ModelXRay class. This provides a boxplot to summarize the extent to which probabilities for given cases tend to change when each variable is adjusted within its empirical range. We'll start here with the Feature Effect Summary for the Random Forest Classifier (Figure 1):




Figure 1: Random Forest Feature Effect Summary

Results in the boxplots are sorted by median effect in changing probabilities. The Random Forest classifier is highly sensitive to changes in a few of the variables, with household income and number of vehicles being the most important. These results make intuitive sense, particularly the idea that income is the variable that models should be most sensitive to. There is also some variability in the sensitivity of the model to key variables. Changes in most variables appear to make little difference to results. This includes whether the household has children or whether it is a single parent household.

The Decision Tree summary (Figure 2) shows sensitivity to a similar set of variables, but with much less consistency; the boxplots are very wide, meaning that there is large variability in the extent to which changes in a given variable effect model results. This makes sense in comparison to the Random Forests results, given that Random Forests takes an average of results from Decision Tree models, reducing variability. Results for Logistic Regression (Figure 3) are dominated by household income to a greater extent than the tree-based results, and shows consistency in sensitivity. Strangely, this model is not sensitive to number of vehicles, which is the second most important variable for the two decision tree related models (and one that I expected to be important). Finally, the nearest neighbors classifier shows almost no sensitivity to most variables other than household income.

The Random Forests model showed sensitivity to different variables in a pattern that is most similar to expectations, but had substantial variability. Logistic regression provided consistent results, but was not sensitive to vehicle ownership, which was surprising.




Figure 2: Decision Tree Feature Effect Summary




Figure 3: Logistic Regression Feature Effect Summary




Figure 4: K Nearest Neighbors Feature Effect Summary

## Evaluation: Interpretability

The final criterion is interpretability. Logistic regression/logit models are a good starting place for this discussion; as linear regression models, their formulations are familiar and easily interpretable. For our fitted logistic regression classifier in scikit-learn, we can output estimated coefficients and format them as a familiar regression equation or random utility equation as below. This allows us to, for instance, calculate odds ratios or tradeoffs between different variables, and gives us a sense of the importance of different variables (given the units of variables are known).

```
Utility = B_0 + -0.137489 * fs_recode + -0.197245 * laptop_recode + 0.011107 * VEH +
-0.143870 * english_hh + -0.131076 * single_parent + -0.310259 * nonfamily +
0.000006 * HINCP + -0.004984 * HUGCL + -0.103327 * children +
-0.301032 * good_english_speaker + -0.367452 * multigen +
-0.190938 * NR + -0.081423 * unmarried_partner + 0.000830 * samesex_marriage
```

Machine learning models are often thought of as black-box models because of the difficulty of interpreting them. Decision tree models are sometimes thought of this way, though visualization techniques can help illuminate what's going on under the hood of these models. The scikit-learn library includes functionality to output decision tree rules as a GraphViz dot file, which can be converted to an image. This provides a visual tree that can help understand how results break down. A small section of our trained Decision Tree classifier is visualized in Figure 5.

A Random Forests model can be visualized in a similar way, using one sub-estimator at a time, but this is difficult to decipher and does not provide a quick way to visually or intuitively interpret the model. The interpretability of both Decision Trees and Random Forests is also subject to the tree depth as well: the deeper the tree, the more difficult to interpret. Finally, nearest neighbor models are also difficult to interpret and provide no quick, intuitive visual summary of model structure, at least beyond very small dimensions. In some ways, while the idea of nearest neighbors is very simple, it is the least interpretable model I evaluated.

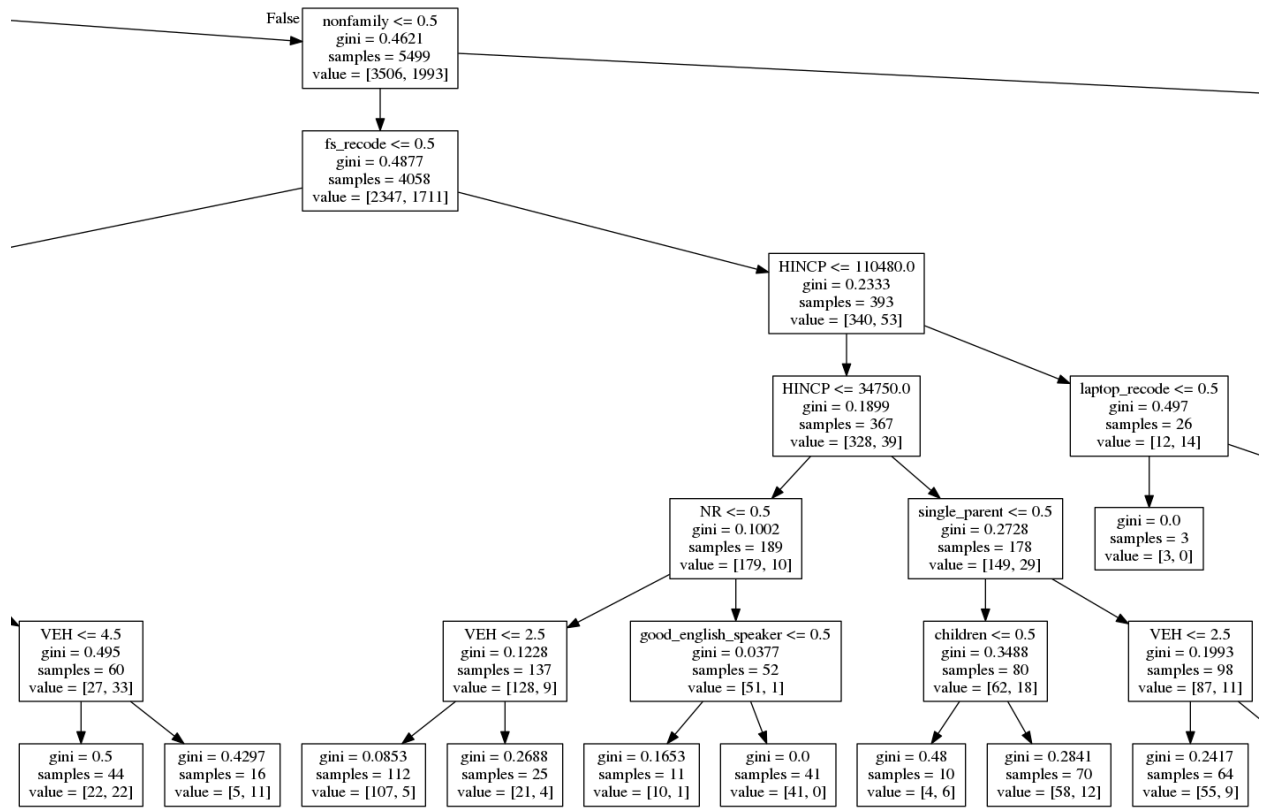


Figure 5: Subset of visualized decision tree



## Conclusion

Decision Trees and their ensemble version, Random Forests, provide a potentially promising alternative to logit modeling for classification problems in the UrbanSim model ecosystem. They are more accurate in their predictions than logistic regression, at least on the California PUMS dataset, if by only a small percentage. Random Forest models also provide the sensitivity to different variables that is closest to my personal expectation, but fall short in their interpretability. This analysis is only a start, though, and these three evaluations should be revisited in a more thorough analysis. There are other major improvements that could be made to this analysis, including using more sophisticated methods for feature selection, testing more classifiers, and performing a more expansive grid search for hyperparameters. As it stands, interpretability seems to be the least important criterion for implementing an alternative classifier in UrbanSim. Therefore, Random Forests emerges as the best possibility for an alternate method for predicting tenure.